

Zubair Nizami

Senior AI Product Manager · Consumer · Infrastructure · Fintech

zubairnizami.com

zubair.nizami@yahoo.com

SF Bay Area

I ship AI products from zero to revenue in **weeks, not quarters**. Eight years across AWS, Meta, Zoox, and Apple, plus **three live consumer tools built solo**. I work at the intersection of technical depth and product taste: multi-agent orchestration, LLM pipelines, full-stack execution, and GTM strategy in the same week.

\$84M →
\$216M

Revenue growth, AWS
Marketplace

157%

Partner-led GTM revenue lift

3

Live consumer products,
shipped solo

\$2,500+

Avg value surfaced per Card
Scout session

WHERE I GO DEEP

Multi-agent orchestration

LLM pipelines

Search & ranking

Data infrastructure at scale

Full-stack shipping

GTM strategy

EIGHT YEARS BUILDING AT

AWS / **Meta** / **Zoox** / **Apple**

EXPERIENCE

Eight years shipping at **scale**.

aws

AWS

\$84M → \$216M

Senior Partner Manager · AI/ML & Cloud GTM

Drove go-to-market for AI/ML and cloud infrastructure, growing partner-led Marketplace revenue 157% year over year and shaping a roadmap influencing \$100M in new revenue.

M

Meta

\$250M / region

Product Manager · Infrastructure & Data Centers

Owned the AI data center Digital Twin platform, using virtual commissioning and simulation to cut infrastructure cost \$250M per region and lift adoption 60%.

Z

Zoox

0 → 1

Senior Technical Product Manager · Autonomy

Shipped autonomy and operator products across a safety-critical AV stack, including a custom handheld controller (0 to 1) and an ML platform supporting 200+ engineers.

(t)

Apple

4 iPhone programs

Product Design Lead · Hardware

Tech-led four iPhone programs (X, XS, 11, 12) and Vision Pro display test systems, from concept through mass production, improving manufacturing yield 50%.

⚡

Plus, three products built solo

Live in production

Full-stack AI · Next.js, TypeScript, multi-model LLM

Apply Studio, Card Scout, and Neural Mob. Problem, users, and results on the next three pages.

LIVE

applystudio.net

Apply Studio

AI job application platform. From a job post to a tailored resume, cover letter, and answers in under 90 seconds.

<90s per application

20+ jobs in parallel

100+ job boards

THE PROBLEM

A serious job application takes 2 to 4 hours to do well: a keyword-matched resume, a tailored cover letter, and custom answers for every role. So most candidates either apply generically and get filtered out, or spend days on a handful of roles. Multi-agent orchestration makes it possible to parallelize that rewriting across dozens of postings at once, compressing days of work into minutes.

WHO IT'S FOR

EARLY CAREER

The broad applier

Applying widely and needs to stand out without a long track record.

CAREER SWITCHER

The re-framer

Needs existing experience re-positioned for a new domain or function.

SENIOR IC

The selective closer

Applies to fewer roles but wants each one tailored, ATS-ready, and fast.

HOW I SOLVE IT

- An **in-house web scraper** that outperforms third-party tools on hard job sites, extracting descriptions, requirements, and application questions.
- A **multi-agent workflow** across GPT, Claude, and Gemini, with one agent per job running in parallel.
- A **bullet-by-bullet accept, edit, reject UI** so the user keeps their voice while the manual grind disappears.
- **Real-time ATS scoring** and keyword-gap analysis across 20+ jobs before submission. Built and shipped solo in under 6 weeks.

RESULTS

Per-application time cut from ~3 hours to under 90 seconds.

20+ tailored applications generated in under one hour.

ATS keyword coverage surfaced before every submission.

Next.js

TypeScript

Vercel

GPT-5.1

Claude

Gemini

Playwright

LIVE

thecardscout.app

Card Scout

AI credit-card optimizer. Real math on your real spending, not affiliate rankings.

\$2,500+ avg annual value

100+ cards analyzed

2 min vs hours

THE PROBLEM

The average household leaves \$1,500 to \$3,000 in credit-card rewards unclaimed each year. Most comparison sites are revenue-optimized for banks, not users, surfacing the highest affiliate payout rather than the highest actual value for a given spending profile. Combining a structured valuation engine with multi-model reasoning personalizes recommendations at a level no static site can match.

WHO IT'S FOR (three entry points, mapped to expertise)

BEGINNER

"I'm starting out"

No cards yet. Answers 3 questions and gets ranked matches for goals and spend.

INTERMEDIATE

"I have cards already"

Wants to see the gaps in a current wallet and exactly what to add next.

ADVANCED

"I know what I want"

Browses and filters 100+ cards by issuer, rewards type, and annual fee.

HOW I SOLVE IT

- Treat rewards as a **valuation problem, not a search problem**, modeling expected annual value across 100+ cards per user profile.
- A **TypeScript valuation engine** computes expected value per card and per card combination from spend and category multipliers.
- **Bank-specific crawlers** keep card data, APRs, and bonus structures fresh automatically, so nothing goes stale.
- **Multi-model orchestration** (GPT, Claude, DeepSeek) for reasoning and explanations, with transparent ranked results and full earnings breakdowns.

RESULTS

Average session surfaces \$2,500+ in net annual card value.

Research time cut from hours to under 2 minutes.

100+ cards analyzed per session, with the "why" shown for every rank.

Next.js

TypeScript

Postgres

GPT-5

Claude

DeepSeek

Playwright

LIVE

neuralmob.xyz

Neural Mob

Ask once, three minds answer. A multi-model debate engine with a judge and visible reasoning.

3 models debate

1 scoring judge

100% reasoning visible

THE PROBLEM

A single model optimizes for plausibility, not correctness. There is no forcing function and no adversarial pressure, so you get confident answers that can be subtly wrong. Multi-agent debate is a known research technique but is rarely productized with user-visible transparency. The tooling now exists to route a query to several models in parallel, collect structured outputs, and score reasoning quality at production speed.

WHO IT'S FOR

BUILDERS & PMS

The model comparer

Wants to see how GPT, Claude, and Gemini differ before trusting one for a task.

DECISION-MAKERS

The high-stakes asker

Needs a higher-confidence answer on an ambiguous, consequential question.

AI POWER USERS

The skeptic

Does not trust one model and wants a judged, fully transparent consensus.

HOW I SOLVE IT

- **Agent roles, system prompts, and model assignments** defined for parallel debate routing, so each agent brings distinct priors.
- A **scoring judge** evaluates every response against a rubric that penalizes hedging and rewards specificity and evidence.
- **Full attribution and reasoning traces** on the result surface, so users see exactly why the winning answer won.
- **Graceful degradation** when an agent times out or returns malformed output. This is core AI PM craft: agent design, failure modes, scoring non-deterministic outputs, and building trust.

RESULTS

Adversarial debate surfaces stronger answers than any single model.

Every answer ships with a traceable scoring rationale.

Built and deployed solo. Live product with real users.

Next.js

TypeScript

Vercel

GPT-5.1

Claude

Gemini

FOR RECRUITERS

Let's build AI products **that ship.**

Target role	Senior or Staff AI Product Manager, where AI is the product, not a feature bolt-on.
Best-fit areas	LLM orchestration · AI-native consumer products · Search & ranking · Data infrastructure at scale
Strengths	Multi-agent orchestration, full-stack shipping, technical depth, and GTM, in the same week.
Location	SF Bay Area · Open to remote · Available within 2 to 4 weeks
Now shipping	Apply Studio v2 · Interview Coach · Agent memory systems

Ready to build AI products that ship?

Let's talk.

zubair.nizami@yahoo.com · zubairnizami.com